

- Lecturer is Jonas.Skeppstedt@cs.lth.se with office E:2190
- Course site is <https://cs.lth.se/edag01>
- Simpler alternative is <https://jonasskeppstedt.net>  
but the Discord server and Tresorit directory are more useful
- You find course files by following the link *here* at the course home page
- You will get an account on an IBM POWER8 Linux server
- You can work on other machines if you wish but performance measurements are to be done on this.

# Lab computer: IBM S882L

- POWER8
- 1 TB RAM
- 20 cores
- 160 hardware threads
- Cost around 1500 EUR from ebay
- `ssh -Y stilid@power8.cs.lth.se`

# Contents of the course

- F1 Introduction to C
- F2 Labs and project: linear and integer programming
- F3 More C
- F4 Instruction set architectures: POWER
- F5 Types, conversions, and linkage
- F6 Superscalar processors: POWER8
- F7 Declarations and expressions
- F8 Cache memories
- F9 Statements and the C preprocessor
- F10 Performance analysis
- F11 The C Standard library
- F12 Optimizing compilers

Sedgewick and Flajolet in "An Introduction to the Analysis of Algorithms":

*The quality of the implementation and properties of compilers, machine architecture, and other major facets of the programming environment have dramatic effects on performance.*

*You will learn the C language in detail and a methodology to maximize algorithm performance on a modern computer*

*To write efficient code, you need competence in:*

- Mathematics, algorithms and data structures
- The C programming language and UNIX C programming tools
- Pipelined and superscalar processors
- Cache memories
- What optimizing compilers can do for you — and what you need to fix yourself

# Contents Lecture 1

- The purpose of learning C
- Some simple C programs

# Some views of C

- The *other* language for high-performance, FORTRAN, is mainly focussed on numerical computing and not for writing code eg for embedded systems, operating system kernels, or compilers.
- Very often other languages such as Clojure, Python, Scala, Haskell, Lisp, Prolog, Ada, Java, C++, Mathematica, or Matlab are preferable because they have many convenient features which enable faster program development.
- When performance in terms of memory usage and/or speed is *the* most important aspect, however, the programmer must have complete control over what is happening and then the overhead of many language features can lead to inferior performance.

# Your lecturer's relationship with C

- C is great but not ideal for *everything*.
- It is my favorite language since 1988. Just like Lisp and Prolog, it's nice because it's beautiful, powerful, and is simple.
- I have written the second ISO validated C99 compiler, after edg.com.
- If I would manage a large software project with several million lines of code, I would use C.
- I will not try to convince you that C "is best" because there is no such thing as a best language.

# Principles of the C Programming Language

- Trust the programmer
- Don't prevent the programmer from doing what needs to be done
- Keep the language small and simple
- Provide only one way to do an operation
- Make it fast, even if it is not guaranteed to be portable
- Support international programming

# Writing a C program

```
#include <stdio.h>

int main(int argc, char** argv)
{
    printf("hello, world\n");
    return 0;
}
```

- A Java methods is called a function in C.
- A C program must have a **main** function.
- A C function must be declared before it is used.

# The C preprocessor

- The command `#include <stdio.h>` reads a file with a declaration of `printf`.
- Commands in a C file which start with a hash, `#`, are performed by the C preprocessor before the compiler starts.
- You can run the preprocessor by typing `cpp`.
- The preprocessor can include files and deal with macros, eg `INT_MAX` is the largest number of type `int`.
- Notice that `cpp` knows nothing about C syntax.

# Installing the gcc and clang compilers on Windows

- Install Windows subsystem for Linux
- See Tresorit and the file links.txt for links to youtube videos (and in the comments part of this video)
- Click on the Ubuntu app and you will get a terminal window.
- Become Ubuntu administrator by typing and press return (or enter)  
`sudo bash`
- Update some files by typing:  
`apt update`
- and then  
`apt upgrade`
- and install  
`apt install gcc clang`
- and to leave administrator mode type  
`exit`

# Installing the clang compiler on a Mac

- Search for and open a terminal window on your Mac
- Then type

```
xcode-select --install
```

- Other compilers can be installed using the brew system but you don't need to use them

# Installing Linux on a computer with Apple silicon

- <https://asahilinux.org/>
- Copy the command there and paste into a terminal
- You need to be administrator of your computer
- This way you can use Valgrind on your Mac laptop
- If you make serious mistake, you can destroy your computer
- HDMI does not work yet
- Tresorit app is only for macOS, X86 Linux and Windows

# Compiling a C program

- In this course we will use the GNU C compiler, called `gcc`.
- To compile one or more C files to make an executable program type `gcc hello.c`
- The command `gcc` will first run `cpp`, then the C compiler, and then two more programs called an assembler and a link-editor.
- Later in the course you will learn about assembler and the operating system course you can learn about link-editors.
- For this course, `gcc` takes care of the link-editor and tells it to produce an executable file.

# Running a C program

- By default the executable file (made by typing `gcc hello.c`) is called **a.out**.
- To execute it in Linux (or MacOS X, or another UNIX), type `./a.out`.
- You can tell gcc that you want a certain name: `gcc hello.c -o hello`.
- Now you type `./hello`.

# Separate compilation

- If you have many big source code files, it is a waste of time to recompile all files every time.
- You can tell gcc to compile a file and produce a so called object file (has nothing to do with object-oriented programming).
- `gcc -c hello.c`
- `gcc hello.o`
- The above two lines are identical to `gcc hello.c` but useful if you have many files. The second line should then contain all .o files.

# Example of I/O: scanf and printf

```
#include <stdio.h>
int main(int argc, char** argv) // same as String args[]
{
    int      a;
    float    b;
    double   c;

    scanf("%d %f %lf", &a, &b, &c);
    printf("%lf\n", a + b + c);
}
```

- %d for int, %f for float, and %lf for double.
- The program will read three numbers from input and print the sum.

# More about the previous example

- In the call to the function `scanf`, we need `&` to tell the compiler that the variables should be modified by the called function.
- This does not exist in Java. You cannot ask another method to modify a number passed as a parameter to the method.
- Other useful format-specifiers include:
  - `%x` for a hexnumber (base 16),
  - `%s` for a string,
  - `%c` for a char,

# Writing to files in C

```
#include <stdio.h>
int main(int argc, char** argv)
{
    int      a = 1;
    float    b = 2;
    double   c = 3;
    FILE*   fp;

    fp = fopen("data.txt", "w"); // open the file for writing.
    fprintf(fp, "%d %f %lf\n", a, b, c);
    fclose(fp);
}
```

- This will create a new file on your hard disk.

# Reading from files in C

```
#include <stdio.h>
int main(int argc, char** argv)
{
    int      a;
    float    b;
    double   c;
    FILE*   fp;

    fp = fopen("data.txt", "r"); // open the file for reading.
    fscanf(fp, "%d %f %lf", &a, &b, &c);
    fclose(fp);
}
```

- Note again the & since fscanf will modify the variables.

# Three ways to make arrays in C

```
#include <stdio.h>
#include <stdlib.h>

int size = 10;

int main(int argc, char** argv)
{
    int      a[10], n, i;
    int*    b;
    int      c[size];           // called a variable length array.

    sscanf(argv[1], "%d", &n); // assumes program is run eg as $ ./a.out 10

    b = calloc(n, sizeof(int)); // like Java's b = new int[n];

    for (i = 0; i < n; i += 1)
        b[i] = i; // use b as if it was an array

    free(b);
}
```

# Explanation of the previous slide

- The **a** and **c** arrays are allocated with other local variables.
- Note that **a** and **c** are "real" arrays.
- On the other hand, **b** is like an array in Java for which you must allocate memory yourself. Use **new** in Java and eg **calloc** in C.
- Java automatically takes care of deallocating the memory of objects.
- In C you must do it yourself using **free**.
- The variable **b** is not an array — it is a pointer.

# Variable length array in C99 and C11

```
int fun(int m, int n)
{
    int      a[n];
    int      b[m][n];
}
```

- Before C99 the above was illegal due to m and n are not constants.
- In C99 it is OK to write like that but only for local variables.
- Most C compilers still only support C89 and thus it may be wise to stick to that at least sometimes.
- Variable lengths arrays are only optional in C11.

# Class in Java vs Struct in C 1(4)

- C has no classes.
- C has structs which are Java classes with everything public and no methods.

```
struct s { // this s is a tag.  
           int      a;  
           int      b;  
} s;       // this s is a variable identifier.
```

- Struct names have a so called **tag** which is a different namespace than variables and functions: so the above declares a **struct s** which is a type and a variable **s**.
- If we write **Link p** in Java we declare **p** to be a reference but not the object itself whereas **s** above is the *real* object, or data.

## Class in Java vs Struct in C 2(4)

- In Java we can declare a List class something like this:

```
class List {  
    List    next; // Next is a reference to another object.  
    int     a;  
    int     b;  
}
```

- **next** above only holds the address of another object but **next is not a List object itself**. The list does not contain a list.
- Java let's you use pointers conveniently without giving you too much head ache.
- C does not.

# Class in Java vs Struct in C 3(4)

- We cannot write the following in C:

```
struct list_t {  
    struct list_t      next;    // Compilation error!!  
    int                a;  
    int                b;  
};
```

- It is impossible to allocate a list within the list!
- We really want to declare **next** to simply hold the address of a list object.
- In C this is done as: **struct list\_t\* next;** which makes **next** a pointer.

# Class in Java vs Struct in C 4(4)

- The following is correct in C:

```
struct list_t {  
    struct list_t*      next;  
    int                a;  
    int                b;  
};
```

- After going into pointers in more detail we will see how to avoid typing **struct list\_t** more than twice using **typedef**.

- As you all know, your computer has something called **memory**.
- It is sufficient to view it as a huge array: **char memory[4294967296];**
- It is preferable in the beginning to view it as: **int memory[1073741824];**
- Forget about strings for the moment. Now our world consists only of ints.
- As you know, a compiler translates a computer program into some kind of language which can be understood by a machine.
- That has happened for the software in everybody's mobile phone.

# Instructions

- You will see more details about it later, but the C program which controls your phone is translated to commands which are numbers and can be represented as ints.
- These ints are also put in the memory.
- We can for instance put the instructions at the beginning of the array.
- The instructions will occupy a large number of array elements.
- No problem — our array is huge.

# Global variables in memory

```
int x = 12;

int main()
{
    return x * 2;
}
```

- We also put the variable `x` in the memory.
- This program will have a few instructions for reading `x` from memory, multiplying with two, and returning the result.
- It is a good idea to put `x` after the instructions: next page

# Memory layout

0	READ from 3 into R	read the data in $x$ from memory at address 3
1	MUL 2	$R = R * 2$
2	RETURN	return $R$
3	12	$x$ lives here

- The array element where we have put a variable is called its **address**
- The instructions above are not written as integers but rather as commands to make them more readable.
- An instruction is represented in memory as a number however.
- It would be too complicated to demand that the hardware should read text such as **MUL** — it is easier to build hardware if there simply is a number which means multiplication.

# Function calls and local variables

- When you call a function or method, all the local variables must be stored somewhere.
- It is a convention to put them at the end of the memory array.
- The local variables of the main function are put at the very end of the array.
- When main calls a function, its local variables are put just before main's.
- In general, when a new function starts running, it puts its local variables at the last (highest index) unused memory array elements.
- This works like a stack of plates: main is at the bottom and you put newly called functions on the plate at the top.

# The Stack

```
int main()          int f(int a)          int g(int a)
{
    int x = 12;      {
                      int b = a+1;      {
    return f(x);    }          return g(b+2); }      return a + 3;
}
```

1073741817	15	a in g lives here.
1073741818		return address from g is here.
1073741819	13	b in f lives here.
1073741820	12	a in f lives here.
1073741821		return address from f is here.
1073741822	12	x in main lives here.
1073741823		return address from main is here.

- When a function returns, it deallocates its memory space.
- This is managed by the compiler which uses a register for holding the current free memory index, called the **stack pointer**.

# Pointers

```
int x = 12;
int *p;
int main()
{
    p = &x;
    *p = 13;
    return x * 2;
}
```

- A pointer is just a variable and it can hold the address of another variable.
- When **p** points to **x**, writing **\*p** accesses **x**.

# Memory layout

	instruction/data	Java	comment
0	STORE 6 at 7	MEMORY[7] = 6	&x is put in element 7, ie p
1	READ from 7 into R	R = MEMORY[7]	read data in p: R=6
2	STORE 13 at R	MEMORY[R] = 13	*p = 13
3	READ 6 into R	R = MEMORY[6]	fetch the value of x
4	MUL 2	R = R * 2	multiply x and R
5	RETURN	return R	
6	12		x lives here
7	0		p lives here

# More about pointers

- In Java, you have used pointers all the time, but they are called object references.
- Suppose you have **Link p**, then **p** is a pointer.
- In Java, pointers can only point at objects.
- The address of some object is, as you might know, the location in memory where that object lives, ie just an integer number.
- In Java, **new** returns the address of a newly created object.
- In C, **new** does not exist and instead a normal function is used (malloc or calloc).

# More about pointers

- In C, but not in Java, the programmer can ask for the address of almost anything and thus get a pointer to that object (or function).
- To change the value of a variable in a function, you need to pass the address of the variable as a parameter to the function:

```
void f(int* a)
{
    *a = 12;
}
```

```
void g()
{
    int      b;
    f(&b);
}
```

# More about pointers

- If the type of the variable is a pointer, then you will need two stars:

```
void f(int** a)
{
    *a = NULL;
}
```

```
void g()
{
    int*     b;
    f(&b);
}
```

# More about pointers

- To return multiple values in Java, you create and return an extra object.
- Option 1 in C: use a plain struct which is allocated on the stack.
- Option 2 in C: Pass additional arguments as pointers (preferable).

```
struct s f()
{
    struct s a;
    a.x = 1;
    a.y = 2;
    a.u = 3;
    return a;
}
```

```
void g(int* x, int* y, int* u)
{
    *x = 1;
    *y = 2;
    *u = 3;
}
```

# Arrays vs Pointers

- Arrays and pointers are not equivalent!
- An array declares storage for a number of elements, except when it is a function parameter:

```
int fun(int a[], int b[12], int c[3][4]);  
int fun(int *a, int *b,     int (*c)[4]);  
int main()  
{  
    int      x, y[12], z[4];  
    fun(&x, y, &z); // valid.  
}
```

- The compiler changes the first [ ] to \* for array parameters.
- Array parameters are not arrays. They are pointers.
- Doing so avoids copying large arrays in function calls.

# C has row-major matrix memory layout

```
int c[3][4] = { { 1, 2, 3, 4}, { 5, 6 }, { 7 } };  
int i, j;  
for (i = 0; i < 3; i++)  
    for (j = 0; j < 4; j++)  
        x += c[i][j];
```

- In a two-dimensional array, one row is laid out in memory at a time, ie row-major.
- Could also be called "rightmost index varies fastest".

# Arrays as parameters

```
int fun(int c[3][4])
{
    printf("%zu %zu\n", sizeof c, sizeof c[0]);
}
```

- If the output is "8 16", what conclusions can we draw about the size of a pointer and the size of an int?
- Answer: a pointer is eight bytes and an int is four bytes.
- The variable c in the function is simply a pointer: `int (*c)[4]`.

# Representation of array references

- $a[i]$  is represented as  $*(a+i)$  internally in the compiler.

```
int main()
{
    int      a[10], *p, i = 3;

    /* the following are equivalent: */

    i[a];
    a[i];
    p = a; p[i]; i[p];
    p = a+i; 0[p]; p[0]; *p;
}
```

# Memory allocation in C

- ① Variables with static storage duration (globals, static).
- ② Stack variables.
- ③ `alloca(size_t size)` takes memory from the stack.
- ④ `malloc/calloc/realloc` take memory from the heap.

# Use tools to find memory errors

- Memory errors:
  - Use pointer which does not point to anything
  - Index out of bounds
  - Forget to free — called a memory leak
  - Free twice
- Two tools you will use in Lab 3
  - Valgrind
  - Google Sanitizer

# Global variables and functions

- Visible from others source files.
- Automatically set to zero unless there is an initializer:

```
int x;  
int y = 1;  
int f() { return x * y; }
```

- Often it is best to avoid global variables due to:
  - Compilers are not good at using them efficiently
  - They sometimes make it more difficult to understand the program

# Static variables and functions

- Similar to global variables and functions

```
static int x;  
static int y = 1;  
static int f() { return x * y;}
```

- Only visible in the scope it is defined
- Functions can only be defined at file scope — no nested functions!
- Always use static instead of global unless the symbol is "exported" to other files
- There is no syntax in C to export symbols — use a header file with declarations

# Stack variables

- Easy for compilers to use efficiently
- Don't use huge arrays since the stack may be too small
- You can use a struct as a parameter and return value — but not array
- As we saw arrays are converted to pointers in the declaration
- There is no syntax to return an array — only a pointer:

```
int a[10];  
int* f()  
{  
    return a; // ok  
}
```

```
int* g()  
{  
    int a[10];  
    return a; // bad idea  
}
```

- The pointer returned from g becomes invalid immediately

# Stack variable

- No automatic initial value — just garbage
- We can initialize a struct or array:

```
int main()
{
    int      a[10] = { 1, 2, 3 };

}
```

- Zero is used for the "missing" expressions

# alloca

- Takes memory from the stack
- Automatically deallocated at function return
- Problem 1: **alloca** is not standard.
- Problem 2: if no memory is available, **NULL** is not returned (as for **malloc/calloc**).
- Somewhat bad reputation, but nevertheless used.
- Much more efficient than **malloc/calloc**.

# Heap memory

- `void* malloc(size_t s);`
- `void* calloc(size_t n, size_t s);`
- `void* realloc(void* p, size_t s);`
- `void free(void* p);`

- Using Java new or malloc/free takes time
- Sometimes a free-list is useful
- Instead of calling free, put it aside for future use
- Instead of calling malloc, check if there already is something put aside
- With "put aside" is meant putting it in a list — but don't allocate memory for the list!
- Use the object type itself somehow

# sizeof

- Use the **sizeof** operator when requesting memory.
- The **sizeof** operator either takes a type or an expression as operand:

```
int* p; /* lots of code... */ p = malloc(n, sizeof(int));  
int* q; /* lots of code... */ q = malloc(n, sizeof *q);
```
- The latter is safer: what happens if somebody changes from `int` to `long` and forgets the `sizeof`-operand?